

Cluster Analysis: Unsupervised Learning via Supervised Learning with a Non-convex Penalty

Wei Pan

WEIP@BIOSTAT.UMN.EDU

*Division of Biostatistics
University of Minnesota
Minneapolis, MN 55455, USA*

Xiaotong Shen

XSHEN@UMN.EDU

*School of Statistics
University of Minnesota
Minneapolis, MN 55455, USA*

Binghui Liu

LIUBH024@GMAIL.COM

*Division of Biostatistics and School of Statistics
University of Minnesota
Minneapolis, MN 55455, USA*

Editor: Francis Bach

Abstract

Clustering analysis is widely used in many fields. Traditionally clustering is regarded as unsupervised learning for its lack of a class label or a quantitative response variable, which in contrast is present in supervised learning such as classification and regression. Here we formulate clustering as penalized regression with grouping pursuit. In addition to the novel use of a non-convex group penalty and its associated unique operating characteristics in the proposed clustering method, a main advantage of this formulation is its allowing borrowing some well established results in classification and regression, such as model selection criteria to select the number of clusters, a difficult problem in clustering analysis. In particular, we propose using the generalized cross-validation (GCV) based on generalized degrees of freedom (GDF) to select the number of clusters. We use a few simple numerical examples to compare our proposed method with some existing approaches, demonstrating our method's promising performance.

Keywords: generalized degrees of freedom, grouping, K-means clustering, Lasso, penalized regression, truncated Lasso penalty (TLP)

1. Introduction

Clustering analysis has been widely used in many fields, for example, for microarray gene expression data (Thalamuthu et al., 2006), mainly for exploratory data analysis or class novelty discovery; see Xu and Wunsch (2005) for an extensive review on the methods and applications. In the absence of a class label, clustering analysis is also called unsupervised learning, as opposed to supervised learning that includes classification and regression. Accordingly, approaches to clustering analysis are typically quite different from supervised learning.

In this paper we adopt a novel framework for clustering analysis by viewing it as a regression problem (Pelckmans et al., 2005; Hocking et al., 2011; Lindsten et al., 2011). We explicitly parametrize each multivariate observation, say x_i , with its own centroid, say μ_i . Clustering analysis

is formulated to identify a small subset of distinct values of these μ_i . Since we have observation-specific and over-parameterized μ_i 's, a key question is how to estimate these parameters. Taking advantage of the recent advance in penalized regression (Tibshirani et al., 2005; Shen and Huang, 2010), we propose a novel non-convex penalty for grouping pursuit that data-adaptively encourages the equality among some unknown subsets of parameter estimates, thus effectively realizing clustering. We call our proposed method as *penalized regression-based clustering* (**PRclust**).

An advantage of regarding clustering as a regression problem is its unification with regression, which in turn provides the opportunity to apply or modify many established results and techniques, such as model selection criteria, in regression to clustering. In particular, a notoriously difficult model selection problem in clustering analysis is to determine the number of clusters; already numerous methods exist with new ones constantly emerging (Tibshirani et al., 2001; Sugar and James, 2003; Wang, 2010). Here we propose the use of generalized cross-validation (GCV) (Golub et al., 1979) that has been widely used for model selection in regression for its solid theoretical foundation, computational efficiency and good empirical performance. However, GCV requires estimating the degrees of freedom (df) or effective number of parameters. In clustering analysis, due to the data-adaptive nature of model searches in finding clusters, it is unclear what is, or how to estimate df. Here we propose using a general method called generalized degrees of freedom (GDF) that was specifically developed in the context of classification and regression to take into account the complex effects of data-adaptive modeling (Ye, 1998; Shen and Ye, 2002). To our knowledge, GDF is mainly studied in the context of regression. Again by formulating clustering as regression, we can adapt the use of GDF to our current context. Although not the main point of this paper, we will show that GDF-based GCV performed well in our numerical examples.

In spite of many advantages of formulating clustering analysis as a penalized regression problem, there are some challenges in its implementation. In particular, with a desired non-smooth and non-convex penalty function, many existing algorithms for penalized regression are not suitable. We develop a novel and efficient computational algorithm that combines the difference of convex (DC) programming (An and Tao, 1997) and a coordinate-wise descent algorithm (Friedman et al., 2007; Wu and Lange, 2008).

Due to some conceptual similarity between our proposed PRclust and the popular K-means clustering, we use the K-means as a benchmark to assess the performance of PRclust. In particular, we show that in some complex situations, for example, in the presence of non-convex clusters, in which the K-means is not suitable, PRclust might perform much better. Hence, complementary to the K-means, PRclust is a potentially useful clustering tool. In addition, we consider a related procedure based on hard thresholding pair-wise distances between observations, called *HTclust*. Although simpler, due to the lack of shrinkage estimation, HTclust may not perform as well as PRclust. Albeit not the focus here, we also propose GDF-based GCV as a general model selection criterion to determine the number of clusters in the above clustering approaches; a comparison with several existing methods demonstrates the promising performance of GCV.

2. Methods

We first present our new method, including its computational algorithm, before comparing it with two related methods. Then we propose a GCV-based method to select the number of clusters.

2.1 New Method: Clustering via Penalized Regression

Given data $X = (x'_1, \dots, x'_n)'$ with $x_i = (x_{i1}, x_{i2}, \dots, x_{ip})'$, we would like to conduct a cluster analysis; that is, we would like to identify group-memberships of the observations such that the within-group similarity and between-group dissimilarity are as strong as possible. There are different ways of defining a similarity/dissimilarity between two observations, leading to various clustering approaches. Here we consider the situation with continuous attributes x_{ik} 's, and define the dissimilarity based on some distance metric, as to be elaborated later. We assume that each data point x_i has its own centroid $\mu_i = (\mu_{i1}, \mu_{i2}, \dots, \mu_{ip})'$, which can be its mean or median (or other measure), depending on the application. Our goal is to estimate μ_i 's while acknowledging the possibility that many μ_i 's would be equal if their corresponding x_i 's are from the same cluster. Hence, we would like to adopt a fused-Lasso-type or fusion penalty (Tibshirani et al., 2005) to encourage the equality of the centroids. In general, we estimate the parameters $\mu = (\mu'_1, \dots, \mu'_n)'$ through minimizing an objective function

$$\hat{\mu} = \arg \min_{\mu} \frac{1}{2} \sum_{i=1}^n L(x_i - \mu_i) + \lambda \sum_{i < j} h(\mu_i - \mu_j),$$

where $L(\cdot)$ is a loss function, for example, the squared error, $h(\cdot)$ is a *grouping* or *fusion* penalty, for example, the L_1 -norm or Lasso penalty (Tibshirani, 1996), and λ is a tuning parameter to be selected. Specifically, with a squared error and Lasso penalty, our objective function is

$$\frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \lambda \sum_{i < j} \|\mu_i - \mu_j\|_1,$$

where $\|\cdot\|_q$ is the L_q -norm. The main idea is that, for the purpose of clustering, we would like to strike a balance between minimizing the distance between the observations and their centroids and reducing the number of centroids via grouping some close centroids together. As pointed out by a reviewer, the general idea with a convex L_q -norm as the fusion penalty has appeared in the literature (Pelckmans et al., 2005; Hocking et al., 2011; Lindsten et al., 2011); here we propose a novel non-convex penalty.

Since it is well known that the Lasso penalty leads to biased parameter estimates (Fan and Li, 2001; Shen et al., 2012), it is more desirable to consider some non-convex penalties; here we propose a new form of the truncated Lasso penalty (TLP) (Shen et al., 2012). For a scalar parameter α and a given tuning parameter τ , TLP is defined as

$$\text{TLP}(\alpha; \tau) = \min(|\alpha|, \tau),$$

which is the L_1 -norm (i.e., Lasso) penalty for a small $\alpha \leq \tau$, but imposes no further penalty for a large $\alpha > \tau$. Importantly, $\text{TLP}(\alpha; \tau)/\tau$ tends to the L_0 -norm of α , $L_0(\alpha) = I(\alpha \neq 0)$, as $\tau \rightarrow 0^+$.

If two observations, x_i and x_j , come from the same cluster, we would have $\mu_i = \mu_j$; that is, all the components of μ_i are equal to that of μ_j . Hence, to more effectively realize $\mu_i = \mu_j$, we use a group penalty that encourages simultaneous equality between all the components of μ_i and μ_j (Yuan and Lin, 2006). Again, to alleviate the bias of the usual convex L_2 -norm (or more generally, L_q -norm for $q > 1$) group penalty, we propose a novel and non-convex group penalty based on TLP, called group TLP or simply gTLP, defined as

$$\text{gTLP}(\mu_i - \mu_j; \tau) = \text{TLP}(\|\mu_i - \mu_j\|_2; \tau).$$

As to be shown, the group TLP performs much better than the Lasso (and other L_q -norms). Note that the group TLP has not been used before.

In this paper, we consider the use of the squared error exclusively, though other loss functions can be used as discussed later. Depending on the use of the penalty, we have two ways to estimate μ_i 's:

$$\begin{aligned}\hat{\mu} &= \arg \min_{\mu} \frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \lambda \sum_{i < j} \|\mu_i - \mu_j\|_1, \\ \hat{\mu} &= \arg \min_{\mu} \frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \lambda \sum_{i < j} \text{TLP}(\|\mu_i - \mu_j\|_2; \tau).\end{aligned}$$

Once we have $\hat{\mu}_i$, then the observations with an equal $\hat{\mu}_i$ are assigned to the same cluster.

2.2 Computing

The above grouping penalties are not separable in μ_i 's in the sense that they cannot be written as a sum of the terms, each of which is a function of a single μ_i only. With the above non-separable penalties, the efficient coordinate-wise algorithm may not converge to a stationary point (Friedman et al., 2007; Wu and Lange, 2008). To develop an efficient coordinate-wise algorithm, we reparametrize by introducing some new parameters and then apply the quadratic penalty method (Nocedal and Wright, 2000). Specifically, we define $\theta_{ij} = \mu_i - \mu_j$ for $1 \leq i < j \leq n$, and then modify the new objective function accordingly as:

$$\begin{aligned}S_L(\mu, \theta) &= \frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \frac{\lambda_1}{2} \sum_{i < j} \|\mu_i - \mu_j - \theta_{ij}\|_2^2 + \lambda_2 \sum_{i < j} \|\theta_{ij}\|_1, \\ S(\mu, \theta) &= \frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \frac{\lambda_1}{2} \sum_{i < j} \|\mu_i - \mu_j - \theta_{ij}\|_2^2 + \\ &\quad \lambda_2 \sum_{i < j} \text{TLP}(\|\theta_{ij}\|_2; \tau).\end{aligned}$$

For $S_L(\mu, \theta)$, the first two terms are quadratic (and thus differentiable and convex) while the third is non-smooth but separable and convex, so the coordinate-wise descent algorithm can be applied and will converge to a global minimum (Tseng, 2001); its updates at iteration $m+1$ are

$$\begin{aligned}\hat{\mu}_i^{(m+1)} &= \frac{x_i + \lambda_1 \sum_{j>i} (\hat{\mu}_j^{(m)} + \hat{\theta}_{ij}^{(m)}) + \lambda_1 \sum_{j<i} (\hat{\mu}_j^{(m+1)} - \hat{\theta}_{ji}^{(m)})}{1 + \lambda_1(n-1)}, \\ \hat{\theta}_{ij}^{(m+1)} &= \text{ST}(\hat{\mu}_i^{(m+1)} - \hat{\mu}_j^{(m+1)}, \lambda_2/\lambda_1),\end{aligned}\tag{1}$$

where $\text{ST}(\alpha, \lambda) = \text{sign}(\alpha)(|\alpha| - \lambda)_+$ is the soft-thresholding rule, and $(a)_+$ takes the positive part of a : it equals to a if $a > 0$, and equals to 0 otherwise. By default any scalar operation on a vector is element-wise.

For $S(\mu, \theta)$, the updating formula for μ_i remains the same as in (1). On the other hand, to deal with the non-convex TLP on θ_{ij} 's, we apply the difference of convex programming technique. We

decompose $S(\mu, \theta)$ into a difference of two convex functions $S_1(\mu, \theta) - S_2(\theta)$:

$$\begin{aligned} S_1(\mu, \theta) &= \frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \frac{\lambda_1}{2} \sum_{i < j} \|\mu_i - \mu_j - \theta_{ij}\|_2^2 + \lambda_2 \sum_{i < j} \|\theta_{ij}\|_2, \\ S_2(\theta) &= \lambda_2 \sum_{i < j} (\|\theta_{ij}\|_2 - \tau)_+. \end{aligned}$$

We then construct a sequence of upper approximations iteratively by replacing $S_2(\theta)$ at iteration $m + 1$ by its piecewise affine minorization

$$S_2^{(m)}(\theta) = S_2(\hat{\theta}^{(m)}) + \lambda_2 \sum_{i < j} (\|\theta_{ij}\|_2 - \|\hat{\theta}_{ij}^{(m)}\|_2) I(\|\hat{\theta}_{ij}^{(m)}\|_2 \geq \tau)$$

at the current estimate $\hat{\theta}^{(m)}$ from iteration m , leading to an upper convex approximating function at iteration $m + 1$:

$$\begin{aligned} S^{(m+1)}(\mu, \theta) &= \frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \frac{\lambda_1}{2} \sum_{i < j} \|\mu_i - \mu_j - \theta_{ij}\|_2^2 + \\ &\quad \lambda_2 \sum_{i < j} \|\theta_{ij}\|_2 I(\|\hat{\theta}_{ij}^{(m)}\|_2 < \tau) + \lambda_2 \tau \sum_{i < j} I(\|\hat{\theta}_{ij}^{(m)}\|_2 \geq \tau). \end{aligned} \quad (2)$$

Applying the (block) coordinate-wise algorithm for the group Lasso (Yuan and Lin, 2006), we have

$$\hat{\theta}_{ij}^{(m+1)} = \begin{cases} \hat{\mu}_i^{(m+1)} - \hat{\mu}_j^{(m+1)}, & \text{if } \|\hat{\theta}_{ij}^{(m)}\|_2 \geq \tau; \\ \left(\|\hat{\mu}_i^{(m+1)} - \hat{\mu}_j^{(m+1)}\|_2 - \frac{\lambda_2}{\lambda_1} \right)_+ \frac{\hat{\mu}_i^{(m+1)} - \hat{\mu}_j^{(m+1)}}{\|\hat{\mu}_i^{(m+1)} - \hat{\mu}_j^{(m+1)}\|_2}, & \text{otherwise.} \end{cases} \quad (3)$$

We summarize below our DC algorithm as **Algorithm 1**:

- STEP 1. (Initialization) Compute an initial estimate $(\hat{\mu}^{(0)}, \hat{\theta}^{(0)})$.
 STEP 2. (Iteration) At iteration $m + 1$, compute $(\hat{\mu}^{(m+1)}, \hat{\theta}^{(m+1)})$ that minimizes (2).
 STEP 3. (Stopping rule) Terminate if $S(\hat{\mu}^{(m+1)}, \hat{\theta}^{(m+1)}) - S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) \geq 0$; otherwise go to Step 2 with $m \leftarrow m + 1$.

We have the following convergence result; its proof is given in an appendix.

Theorem 1 *In Algorithm 1, $S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)})$ decreases strictly in m until it terminates in finite steps; that is, there exists an $m^* < \infty$ with*

$$S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) = S(\hat{\mu}^{(m^*)}, \hat{\theta}^{(m^*)}) \quad \text{for } m \geq m^*.$$

Furthermore, $(\hat{\mu}^{(m^)}, \hat{\theta}^{(m^*)})$ is a local minimizer of $S(\mu, \theta)$.*

In implementing the Step 2 of Algorithm 1, we can repeatedly apply the coordinate-wise updates (1) and (3) to minimize (2). Since the objective function (2) is a sum of a differentiable and convex function and a convex penalty in μ and θ (while $\hat{\theta}^{(m)}$ is a known vector), the coordinate-wise descent algorithm will converge to its minimizer (Tseng, 2001). By Theorem 1, we know that this implementation of Algorithm 1 will converge to a local minimum of $S(\mu, \theta)$. In practice, especially in earlier iterations, one may not want to run the coordinate-wise updates fully until convergence in Step 2 to save computing time.

Note that, due to the use of the quadratic/ridge penalty on $\mu_i - \mu_j - \theta_{ij}$, no matter how large λ_1 is used, we cannot obtain exactly $\mu_i - \mu_j - \theta_{ij} = 0$, though their difference tends to 0 as $\lambda_1 \rightarrow \infty$; on the other hand, the ridge penalty is smooth, and thus facilitates the applicability of the coordinate-wise descent algorithm. To enforce the constraint $\mu_i - \mu_j = \theta_{ij}$ approximately, it is desirable to use a large λ_1 ; however, by (1), we see that a large λ_1 effectively reduces the weight of observation x_i 's contributing to estimating μ_i . We fix $\lambda_1 = 1$ throughout, leaving it to section 4 to discuss an alternative algorithm allowing $\lambda_1 \rightarrow \infty$.

Due to the use of Lasso or TLP on θ_{ij} 's, we can obtain exactly $\hat{\theta}_{ij} = 0$ for a large λ_2 . We form clusters based on $\hat{\theta}_{ij}$'s: for any two observations x_i and x_j , if $\hat{\theta}_{ij} = 0$, they are declared to be in the same cluster. We construct a graph G based on an adjacency matrix $A = (a_{ij})$ with elements $a_{ij} = I(\hat{\theta}_{ij} = 0)$; finding clusters is equivalent to finding connected subcomponents of G . It is possible that other more sophisticated graph-based methods can be used to identify clusters, which we leave as a future topic. By default, PRclust is based on the TLP, not Lasso, unless specified otherwise.

2.3 Comparison with Two Related Methods

Our proposed method is closely related to the K-means method, which can be formulated as finding the centroids, say μ_1, \dots, μ_K , for K clusters, where $K \geq 1$ is a tuning parameter. To find the centroids $\mu = (\mu'_1, \dots, \mu'_K)'$,

$$\hat{\mu} = \arg \min_{\mu} \sum_{i=1}^n \|x_i - \mu_{c(i)}\|_2^2,$$

where $c(i)$ maps observation i to cluster $c(i) \in \{1, 2, \dots, K\}$, one of the K candidate clusters. A typical K-means algorithm starts with some initial estimate of μ , assigns each observation to its nearest centroid/cluster, recalculates each centroid, then repeats the above process until convergence. Depending on the initial estimates, the K-means may converge only to a local minimum, hence multiple starts are often used.

It is clear that both the PRclust and K-means aim to identify centroids by minimizing the total L_2 distance between observations and their corresponding centroids, but they approach and formulate the problem differently. PRclust over-parametrizes the centroid for each observation, then via shrinking parameters by grouping pursuit, finds a fewer number of distinct centroids; the final result depends on the specified tuning parameters λ_2 and τ . In contrast, by specifying K , the number of clusters as the only tuning parameter, the K-means starts with some K initial centroids, then assigns each observation to a cluster before updating the centroid estimates. Hence, PRclust differs from the K-means in that PRclust does not explicitly assign an observation to any cluster; clustering is implicitly done after the convergence of the algorithm.

A simple alternative to PRclust is to apply the hard-thresholding rule to pair-wise distances between observations. Suppose that $D = (d_{ij})$ is a pair-wise distance matrix with $d_{ij} = \|x_i - x_j\|_2$. For any threshold d , we can define an adjacency matrix $A = (a_{ij})$ with $a_{ij} = I(d_{ij} < d)$; as in PRclust, we can define any connected subcomponent based on A as a cluster, resulting in a clustering method called **HTclust**, which is named as a "connected components" algorithm in Ng et al. (2002). As a comparison, PRclust defines its adjacency matrix as $a_{ij} = I(\hat{\theta}_{ij} = 0) = I(\|\hat{\mu}_i - \hat{\mu}_j\|_2 < d_{0,ij})$ for some small and possibly (i, j) -dependent threshold $d_{0,ij} > 0$. HTclust is related to agglomerative hierarchical clustering: the latter also starts with each observation being its own cluster, then sequentially merges two nearest clusters until only one cluster left. Indeed, as shown in an appendix,

HTclust is equivalent to the single-linkage hierarchical clustering, in which the distance between two clusters is defined as the shortest distance between any two observations, one from each of the two clusters. An apparent difference between PRclust and HTclust (and hierarchical clustering) is the lack of shrinkage in parameter estimation in the latter, in contrast to that in the former as shown in (1). In general PRclust behaves differently from HTclust, as to be shown in a few examples.

2.4 Selecting Tuning Parameters or Number of Clusters

A bonus with the regression approach to clustering is the potential application of many existing model selection methods for regression or supervised learning to clustering. Here we propose using generalized cross-validation (GCV) that has been used extensively, for example, in selecting the tuning parameter in ridge regression (Golub et al., 1979). GCV can be regarded as an approximation to leave-one-out cross-validation (CV). Hence, GCV provides an approximately unbiased estimate of the prediction error. In our notation,

$$\text{GCV}(\text{df}) = \frac{\text{RSS}}{(np - \text{df})^2} = \frac{\sum_{i=1}^n \sum_{k=1}^p (x_{ik} - \hat{\mu}_{ik})^2}{(np - \text{df})^2},$$

where df is the degrees of freedom used in estimating μ_i 's. For our problem, a naive treatment is to take $\text{df} = Kp$, the number of unknown parameters in μ_i 's, which however does not take into account the data-adaptive nature in estimating μ_i 's in clustering analysis. As to be shown, the naive estimate of df, $\text{NDF} = Kp$, is in general severely under-biased. A better way is to use the generalized degrees of freedom (GDF) (Ye, 1998). We define GDF as

$$\text{GDF} = \sum_{i=1}^n \sum_{k=1}^p \frac{1}{\sigma^2} \text{cov}(\hat{\mu}_{ik}(X), x_{ik} - \mu_{ik}) = \sum_{i=1}^n \sum_{k=1}^p \lim_{\delta \rightarrow 0} E_{\mu} \left[\frac{\hat{\mu}_{ik}(X + \delta e_{ik}) - \hat{\mu}_{ik}(X)}{\delta} \right],$$

where we write $\hat{\mu}_{ik} = \hat{\mu}_{ik}(X)$ to emphasize that the estimate $\hat{\mu}_{ik}$ depends on the data X being used, and e_{ik} is a vector of length np with all elements 0 except a 1 in position ik . Accordingly, we can use Monte Carlo simulations to estimate GDF in the following way:

Step 1. For $b = 1, \dots, B$, repeat Steps 2-3.

Step 2. Generate $\Delta_b = (\delta_{b,1}, \dots, \delta_{b,np})$ with $\delta_{b,i}$ iid $N(0, \nu)$.

Step 3. Conduct a cluster analysis (in the same way as for the original data X) with data $X + \Delta_b$ to yield an estimate $\hat{\mu}(X + \Delta_b)$.

Step 4. For fixed i and k , regress $\hat{\mu}_{ik}(X + \Delta_b)$ on $\delta_{b,ik}$ with $b = 1, \dots, B$; denote the slope estimate as \hat{h}_{ik} .

Step 5. Repeat Step 4 for each i and k . Then an GDF estimate is $\text{GDF} = \sum_{i=1}^n \sum_{k=1}^p \hat{h}_{ik}$.

We used $B = 100$ in Step 1 throughout. In Step 2, the perturbation size (i.e., standard deviation, SD) ν is chosen to be small, typically with $\nu \in [0.5\sigma, \sigma]$, where a common variance $\sigma^2 = \text{var}(x_{ik})$ is assumed for all attributes. As discussed in Ye (1998) and Shen and Ye (2002), often the GDF estimate is not too sensitive to the choice of ν . In Step 3, we apply the same clustering algorithm (e.g., PRclust or HTclust) with any fixed tuning parameter values as applied to the original data X .

We try with various tuning parameter values, obtaining their corresponding GDFs and thus GCV statistics, then choose the set of the tuning parameters with the minimum GCV statistic.

The above method can be equally applied to the K-means method to select the number of clusters: we just need to apply the K-means with a fixed number of clusters, say K , in Step 3, then use the cluster centroid of observation x_i as its estimated mean $\hat{\mu}_i$; other steps remain the same. Again, we try various values of K , and choose $\hat{K} = K$ that minimizes the corresponding GCV(GDF) statistic.

As a comparison, we also apply the Jump statistic to select the number of clusters for the K-means (Sugar and James, 2003). For K clusters, a distortion (or average within-cluster sum of squares) is defined to be

$$W_K = \sum_{i=1}^n \sum_{k=1}^p (x_{ik} - \hat{\mu}_{ik})^2 / (np),$$

and the Jump statistic is defined as

$$J_K = 1/W_K^{p/2} - 1/W_{K-1}^{p/2},$$

with $1/W_0^{p/2} = 0$. We choose $\hat{K} = \arg \min_K J_K$.

Wang (2010) proposed a consistent estimator for the number of clusters based on clustering stability. It is based on an intuitive idea: with the correct number of clusters, the clustering results should be most stable. The method requires the use of three subsets of data: two are used to build two predictive models for the same clustering algorithm with the same number of clusters, and then the third is used to estimate the clustering stability by comparing the predictive results of the third subset when applied to the two built predictive models. For a given data set, cross-validation is used to repeatedly splitting the data into three (almost equally sized) subsets. Wang (2010) proposed two CV schemes, called CV with voting and CV with averaging. We will simply call the two methods as CV1 and CV2.

3. Numerical Examples

Now we use both simulated data and real data to evaluate the performance of our method and compare it with several other methods.

3.1 Simulation Set-ups

We considered five simulation set-ups, covering a variety of scenarios, as described below.

Case I: two convex clusters in two dimensions (Figure 1a). We consider two somewhat overlapping clusters with the same spherical shape, which is ideal for the K-means. Specifically, we have $n = 100$ observations, 50 from a bivariate Normal distribution $N((0,0)', 0.33I)$ while the other 50 from $N((1,1)', 0.33I)$.

Case II: two non-convex clusters in two dimensions (Figure 1b). In contrast to the previous case favoring the K-means, the second simulation set-up was the opposite. There were 2 clusters as two nested circles (distorted with some added noises), each with 100 observations (see the upper-left panel in Figure 3). Specifically, for cluster 1, we had $x_{i1} = -1 + 2(i-1)/99$, $x_{i2} = s_i \sqrt{1 - x_{i1}^2} + \varepsilon_i$, $s_i = -1$ or 1 with an equal probability, ε_i randomly drawn from $U(-0.1, 0.1)$, for $i = 1, \dots, 100$; for cluster 2, similarly we had $x_{i1} = -2 + 4(i-101)/99$, $x_{i2} = s_i \sqrt{4 - x_{i1}^2} + \varepsilon_i$ for $i = 101, \dots, 200$. This

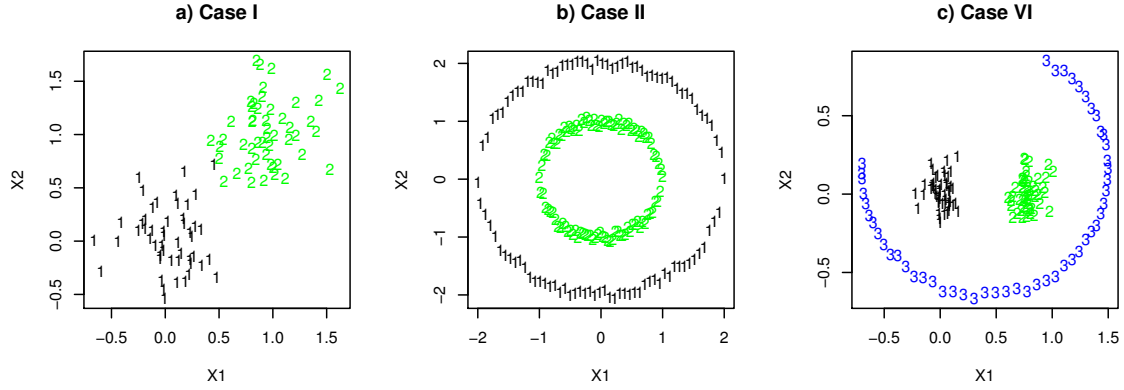


Figure 1: The first simulated data set in a) Case I, b) Case II and c) Case VI.

is similar to the “two-circle” case in Ng et al. (2002), but perhaps more challenging here with larger distances between some points within the same cluster.

Case III: a null case with only a single cluster in 10 dimensions. 200 observations were uniformly distributed over a unit square independently in each of the 10 dimensions. This is scenario (a) in Tibshirani et al. (2001).

Case IV: four clusters in 3 dimensions. The four cluster centers were randomly drawn from $N(0, 5I)$ in each simulation; if any of their distance was less than 1.0, then the simulation was abandoned and re-run. In each cluster, 25 or 50 observations were randomly chosen, each drawn from a normal distribution with mean at the cluster center and the identity covariance matrix. This is scenario (c) in Tibshirani et al. (2001).

Case V: two elongated clusters in 3 dimensions. This is similar to scenario (e) in Tibshirani et al. (2001), but with a much shorter distance between the two clusters. Specifically, cluster one was generated as follows: 100 observations were generated be equally spaced along the main diagonal of a three dimensional cube, then independent normal variate with mean 0 and $SD=0.1$ were added to each coordinate of each of the 100 observations; that is, $x_{ij} = -0.5 + (i-1)/99 + \epsilon_{ij}$, $\epsilon_i \sim N(0, 0.1)$ for $j = 1, 2, 3$ and $i = 1, \dots, 100$. Cluster 2 was generated in the same way, but with a shift of 2, not 10, making it harder than that used in Tibshirani et al. (2001) in each dimension.

Case VI: three clusters in 2 dimension with two spherically shaped clusters inside 3/4 of a perturbed circle (Figure 1c). This is similar to a case in Ng et al. (2002). Specifically, for cluster 1, we generated $x_{i1} = 1.1 \sin(2\pi[30 + 5(i-1)]/360)$ and $x_{i2} = 0.8 \sin(2\pi[30 + 5(i-1)]/360) + \epsilon_i$ for $i = 1, \dots, 50$, where ϵ_i was randomly drawn from $U(-0.025, 0.025)$; 50 observations were drawn from each of the two bivariate Normal distributions, $N((0, 0)', 0.1I)$ and $N((0.8, 0)', 0.1I)$.

For each case, we applied the K-mean, HTclust and PRclust to 100 simulated data sets. For the K-means we used 20 random starts for each $K = 1, 2, \dots, 20$. For HTclust and PRclust, we did grid-searches for d , and (τ, λ_2) respectively. For comparison, we also applied Gaussian mixture-model based clustering as implemented in R package `mclust` (Fraley and Raftery, 2006); for each data set, we fitted each of the 10 models corresponding to 10 different ways of parameterizing the mixture model, for $K = 1, 2, \dots, 20$ clusters, and the final model, including the number of clusters, was selected by the Bayesian Information Criterion (BIC).

Due to the conceptual similarity between our proposed PRclust and spectral clustering (Sclust), we also included the spectral clustering algorithm of Ng et al. (2002) as outlined below. First, calculate an affinity matrix $A = (A_{ij})$ with elements $A_{ij} = \exp(-||x_i - x_j||_2^2/\gamma)$ for any two observations $i \neq j$ and $A_{ii} = 0$, where γ is a scaling parameter to be determined. Second, calculate a diagonal matrix $D = \text{Diag}(D_{11}, \dots, D_{nn})$ with $D_{ii} = \sum_{j=1}^n A_{ij}$. Third, calculate $L = D^{-1/2}AD^{-1/2}$. Fourth, for a specified number of clusters k , we stack the k top eigen-vectors (corresponding to the k largest eigen-values) of L column-wise to form an $n \times k$ matrix, say Z_k ; normalize each row of Z_k to have a unit L_2 -norm. Finally, treating each row of Z_k as an observation, we apply the K-means to form k clusters. There are two tuning parameters γ and k that have to be decided in the algorithm. We used the implementation in the R package `kernlab`, which includes a method of Ng et al. (2002) to select γ automatically; however, one has to specify k . We applied the GCV(GDF) to select k as for the K-means. Unfortunately the function `specc()` in the R package `kernlab` was not numerically stable and sometimes might break down (i.e., exiting with an error message), though it could work in a re-run with a different random seed; the error occurred more frequently with an increasing k . Hence we only considered its use in a few cases by restricting k to 1 to 3.

To evaluate the performance of a clustering algorithm, we used the Rand index (Rand, 1971), adjusted Rand index (Hubert and Arabie, 1985) and Jaccard index (Jaccard, 1912), all measuring the agreement between estimated cluster memberships and the truth. Each index is between 0 and 1 with a higher value indicating a higher agreement.

3.2 Simulation Results

Case I: For the K-means, we chose the number of clusters using Jump, CV1, CV2 and GCV statistics; for comparison, we also fixed the number of clusters around its true value. The results are shown in Table 1. Both the Jump and GCV with the naive $df=np$ methods tended to select a too large number of clusters. In contrast, the GCV(GDF) performed extremely well: it always chose the correct $K = 2$ clusters. Figure 2 shows how GDF and NDF changed with K , the number of clusters in the K-means algorithm, for the first simulated data set. Due to the adaptiveness of the K-means, GDF quickly increased to 150 with $K < 10$ and approached the maximum $df=np = 200$ for $K = 20$. Since GDF was in general much larger than NDF, using GDF penalized more on more complex models (i.e., larger K in the K-means), explaining why GCV(GDF) performed much better than GCV(NDF).

Since the two clusters were formed by observations drawn from two Normal distributions, as expected, the model-based clustering Mclust performed best. In addition, the spectral clustering also worked well.

For PRclust, we searched $\tau \in \{0.1, 0.2, \dots, 1\}$ and $\lambda_2 \in \{0.01, 0.05, 0.1, 0.2, 1\}$. PRclust with GCV(GDF) selecting its tuning parameters performed well too: the average number of clusters is close to the truth $K_0 = 2$; the corresponding clustering results had high degrees of agreement with the truth, as evidenced by the high indices. Table 1 also displays the frequencies of the number of clusters selected by GCV(GDF): for the overwhelming majority (98%), either the correct number of cluster $K_0 = 2$ was selected, or a slightly larger $K = 3$ or 4 with very *high* agreement indices was chosen.

For the first (and a typical) simulated data set, we show how PRclust operated with various values of the tuning parameter λ_2 (while $\lambda_1 = 1$ and $\tau = 0.5$), yielding the solution path for $\hat{\mu}_{i1}$, the first coordinate of $\hat{\mu}_i$ (Figure 3a). Note that, due to the use of a fixed λ_1 , even if all $\hat{\theta}_{ij} = 0$ for

Case	Method	Selection of K	\hat{K}	Rand	aRand	Jaccard
I	K-means	Jump, $K \in [1, 20]$	18.11	0.563	0.119	0.119
		Jump, $K \in [1, 10]$	2.54	0.956	0.911	0.912
		CV1	2	0.984	0.967	0.968
		CV2	2	0.984	0.967	0.968
		GCV, $df=Kp$	29.75	0.536	0.064	0.064
		GCV, $df=GDF$	2	0.984	0.967	0.968
		Fixed	3	0.859	0.717	0.720
		Fixed	4	0.747	0.491	0.495
		Fixed	5	0.704	0.405	0.408
	Mclust	BIC	2.01	0.983	0.966	0.967
	Sclust	GCV, $df=GDF$	2.05	0.973	0.946	0.953
		$K \in [1, 3]$				
	HTclust	GCV, $df=Kp$	60.29	0.524	0.039	0.039
		GCV, $df=GDF$	4.12	0.901	0.802	0.858
		Fixed	2	0.589	0.186	0.583
		Fixed	3	0.711	0.426	0.690
		Fixed	4	0.779	0.562	0.746
		Fixed	5	0.805	0.612	0.760
	PRclust	GCV, $df=Kp$	87.00	0.510	0.009	0.009
		GCV, $df=GDF$	2.35	0.974	0.947	0.953
		Subset, freq=1	1	0.495	0.000	0.495
		Subset, freq=72	2	0.982	0.965	0.966
		Subset, freq=19	3	0.973	0.946	0.946
		Subset, freq=7	4	0.966	0.933	0.933
		Subset, freq=1	5	0.887	0.774	0.788

Table 1: Simulation I results based on 100 simulated data sets with 2 clusters.

a sufficiently large λ_2 , there were still quite some unequal $\hat{\mu}_{i1}$'s, which were all remarkably near their true values 0 or 1. In contrast, with the Lasso penalty, the estimated centroids were always shrunk towards each other, leading to their convergence to the same point at the end and thus much worse performance (Figure 3c). It is also noted that the solution paths with the Lasso penalty were almost linear, compared to the nearly step functions with the gTLP. Figure 3d) shows how HTclust worked. In particular, as pointed out by Ng et al. (2002), HTclust is not robust to outliers: since an “outlier” (lower left corner in Figure 1a) was farthest away from any other observations, it formed its own cluster while all others formed another cluster when the threshold d was chosen to yield two clusters. This example demonstrates different operating characteristics between PRclust and HTclust, offering an explanation of the better performance of PRclust over HTclust.

Case II: Since each cluster was not spherically shaped, and more importantly, the two true cluster centroids completely overlapped with each other, the K-means would not work: it could not distinguish the two clusters. As shown in Table 2, no matter what method was used to choose or fix the number of clusters, the K-means always gave results in low agreement with the truth. The problem with the K-means is its defining a cluster centroid as the mean of the observations assigned




Figure 2: GDF (marked with "G") and NDF (marked with "N") versus the number of clusters, K , in the K-means algorithm for the first simulated data set in Case I. The horizontal line gives the maximum $df=200$.

to the cluster and its assigning a cluster membership of an observation based on its distance to the centroids; since the two clusters share the same centroid in truth, the K-means cannot distinguish the two clusters. Similarly, Mclust did not perform well.

As a comparison, perhaps due to the nature of the local shrinkage in estimating the centroids, PRclust worked much better than the above three methods, as shown in Table 2. Note that, the cluster memberships in PRclust are determined by the estimates of $\theta_{ij} = \mu_i - \mu_j$; due to the use of the ridge penalty with a fixed $\lambda_1 = 1$, we might have $\hat{\theta}_{ij} = 0$ but $\hat{\mu}_i \neq \hat{\mu}_j$.

Since HTclust assigned the cluster-memberships according to the pair-wise distances among the observations, not the nearest distance of an observation to the centroids as done in the K-means, it also performed well.

If the GCV(GDF) was used in Sclust, it would select $\hat{K} = 1$ over $\hat{K} = 2$, even though a specified $\hat{K} = 2$ often led to almost perfect clustering. The reason is that, by symmetry of the two clusters, the two estimated cluster centroids for $\hat{K} = 2$ almost coincided with the estimated centroid of only one cluster, leading to their almost equal RSS (the numerator of the GCV statistics); due to a much larger GDF for $\hat{K} = 2$ than that for $\hat{K} = 1$, the GCV(GDF) statistic for $\hat{K} = 1$ was much smaller than that for $\hat{K} = 2$. Interestingly, an exception happened in four (out of 100) simulations: when Sclust could not correctly distinguish the two true clusters (with low agreement statistics) with $\hat{K} = 2$, it had a smaller GCV(GDF) statistic than that for $\hat{K} = 1$. The results here suggest that, although Sclust may perform well for non-convex clusters with an appropriately chosen γ (as selected by the method

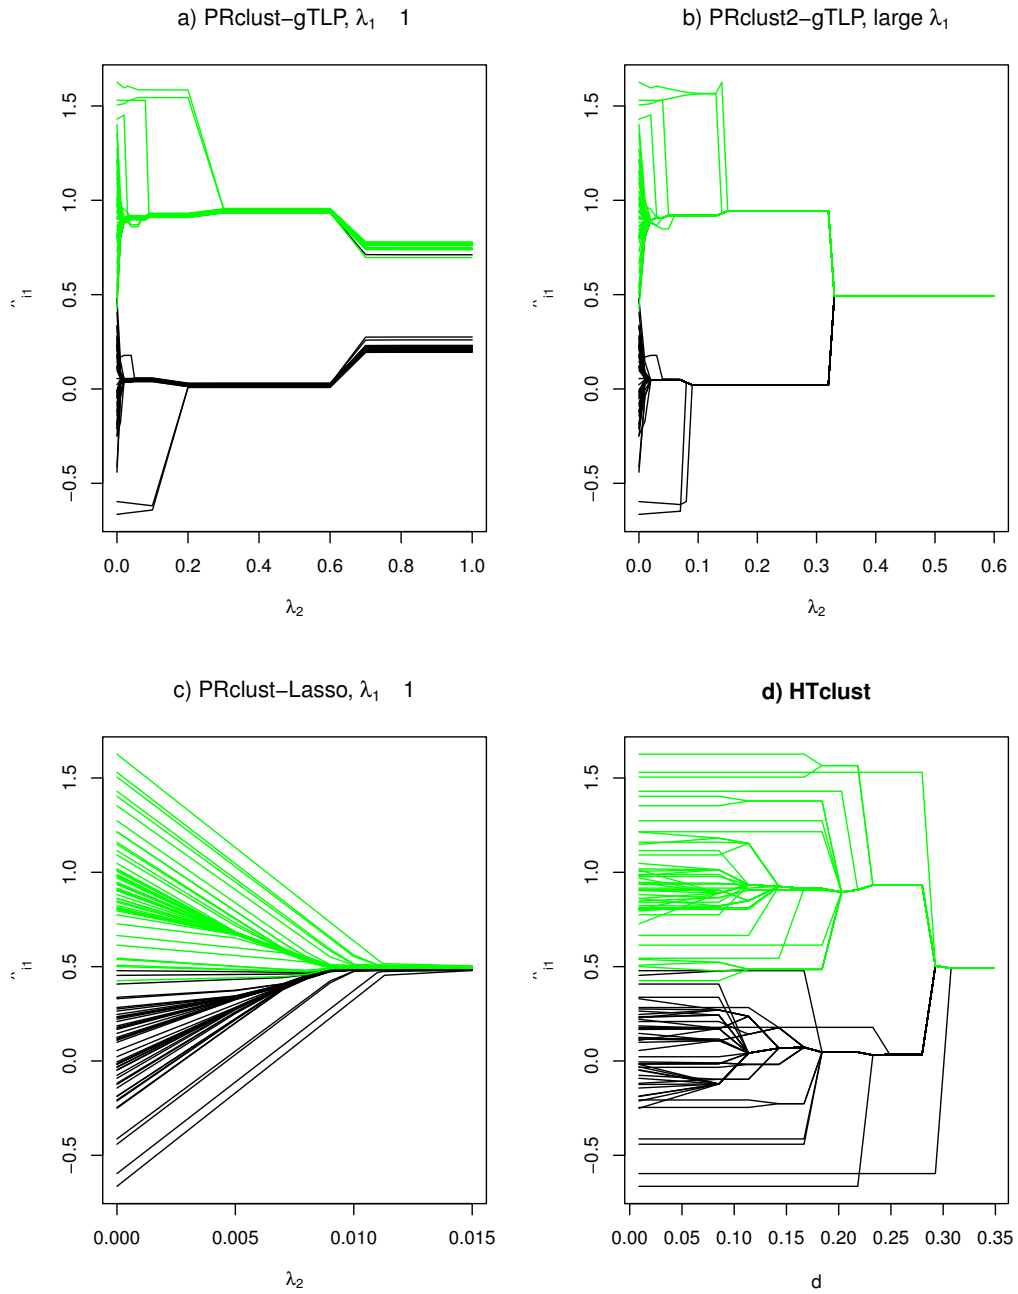


Figure 3: Solution paths of $\hat{\mu}_{i,1}$ for a) PRclust (with gTLP), b) PRclust2, c) PRclust with the Lasso penalty and d) HTclust for the first simulated data set in Case I.

of Ng et al. (2002)), a difficult problem is how to choose the number of clusters; in particular, GCV is not ideal for non-convex clusters.

Case	Method	Selection of K	\hat{K}	Rand	aRand	Jaccard
II	K-means	Jump, $K \in [1, 20]$	18.88	0.557	0.111	0.110
		Jump, $K \in [1, 10]$	9.99	0.597	0.191	0.195
		GCV, df=GDF	2	0.498	-0.005	0.329
		CV1	2	0.498	-0.005	0.329
		CV2	2	0.498	-0.005	0.329
		Fixed	3	0.498	-0.006	0.261
		Fixed	4	0.498	-0.008	0.194
		Fixed	5	0.498	-0.007	0.166
	Mclust	BIC	16.07	0.572	0.141	0.140
	Sclust	GCV, df=GDF	1.06	0.501	0.007	0.493
		$K \in [1, 3]$				
		Subset, freq=95	1	0.497	0.000	0.497
		Subset, freq=4	2	0.498	-0.005	0.329
		Subset, freq=1	3	0.874	0.749	0.748
		Fixed	2	0.980	0.960	0.973
	HTclust	GCV, df=GDF	3.32	0.862	0.724	0.738
		Fixed	2	1.000	1.000	1.000
		Fixed	3	0.881	0.763	0.762
		Fixed	4	0.870	0.739	0.738
		Fixed	5	0.866	0.732	0.731
	PRclust	GCV, df=GDF	2.93	0.895	0.791	0.790
		Subset, freq=21	2	1.000	1.000	1.000
		Subset, freq=66	3	0.880	0.759	0.759
		Subset, freq=12	4	0.810	0.620	0.619
		Subset, freq=1	5	0.746	0.491	0.490

Table 2: Simulation II results based on 100 simulated data sets with 2 clusters.

Cases III-IV: the simulation results are summarized in Table 3. All performed well for the null Case III. Case IV seems to be challenging with partially overlapping spherically shaped clusters of smaller cluster sizes: the number of clusters could be under- or over-selected by various methods. In terms of agreement, overall, as expected, the K-means with GCV(GDF) and Mclust performed best, closely followed by PRclust with GCV(GDF), which performed much better than HTclust.

Cases V-VI: the simulation results are summarized in Table 4. In Case V, all performed perfectly except that the GCV(GDF) over-selected the number of the clusters in the K-means and the two spectral clustering methods. This is interesting since GCV(GDF) seemed to perform well for both HTclust and PRclust. HTclust and PRclust did not yield better clusters than that of the K-means for $K > 2$ clusters, leading to the former two's relatively large GDFs and thus relatively large GCV statistics, while the latter possessed a smaller GDF and GCV, hence GCV(GDF) tended to select a $K > 2$ for the K-means, but not for the other two. Note that the K-means implicitly assumes that all clusters share the same volume and spherical shape, and GCV also implicitly favors such clusters (with smaller within-cluster sum of squares, and thus a smaller GCV statistic). Hence the K-means

Case	Method	Selection of K	\hat{K}	Rand	aRand	Jaccard
III	K-means	GCV, df=GDF	1.00	1.000	1.000	1.000
	Mclust	BIC	1.00	1.000	1.000	1.000
	Sclust	GCV, df=GDF $K \in [1, 3]$	1.00	1.000	1.000	1.000
	HTclust	GCV, df=GDF	1.00	1.000	1.000	1.000
	PRclust	GCV, df=GDF	1.00	1.000	1.000	1.000
IV	K-means	GCV, df=GDF	3.48	0.880	0.748	0.728
		CV1	3.10	0.789	0.575	0.581
		CV2	4.22	0.790	0.558	0.561
	Mclust	BIC	3.50	0.883	0.753	0.732
	HTclust	GCV, df=GDF	6.49	0.589	0.352	0.452
	PRclust	GCV, df=GDF	4.75	0.790	0.612	0.628

Table 3: Simulation Cases III-IV results based on 100 simulated data sets with 1 and 4 clusters, respectively.

Case	Method	Selection of K	\hat{K}	Rand	aRand	Jaccard
V	K-means	GCV, df=GDF	7.03	0.646	0.289	0.288
		CV1	2.00	1.000	1.000	1.000
		CV2	2.00	1.000	1.000	1.000
	Mclust	BIC	2.04	0.995	0.990	0.990
	HTclust	GCV, df=GDF	2.00	1.000	1.000	1.000
VI	PRclust	GCV, df=GDF	2.00	1.000	1.000	1.000
	K-means	GCV, df=GDF	7.95	0.902	0.761	0.704
		CV1	2.00	0.722	0.444	0.497
		CV2	2.00	0.722	0.444	0.497
	Mclust	BIC	8.04	0.906	0.769	0.714
	HTclust	GCV, df=GDF	28.02	0.872	0.678	0.611
	PRclust	GCV, df=GDF	3.08	0.997	0.993	0.993

Table 4: Simulation cases V-VI results based on 100 simulated data sets with 2 and 3 clusters, respectively.

divided an elongated cluster into several adjacent spherical clusters, which were then favored by GCV(GDF).

For Case VI, due to the fact of a non-convex cluster, both the K-means with GCV(GDF) and Mclust over-selected the number of clusters, though their agreement statistics were still high. On the other hand, the K-means with CV1 or CV2 and the two spectral clustering methods seemed to under-select the number of clusters, leading to lower agreement statistics. In contrast, PRclust performed much better while HTclust was the worst.

Method	Selection of K	\hat{K}	Truth: 3 cluster			Truth: 2 cluster		
			Rand	aRand	Jaccard	Rand	aRand	Jaccard
K-means	Jump, $K \in [1, 30]$	27	0.710	0.154	0.122	0.490	0.069	0.077
	Jump, $K \in [1, 10]$	10	0.755	0.330	0.281	0.555	0.178	0.194
	Jump, $K \in [1, 5]$	5	0.767	0.420	0.399	0.662	0.362	0.388
	CV1	2	0.776	0.568	0.595	1.000	1.000	1.000
	CV2	2	0.776	0.568	0.595	1.000	1.000	1.000
	GCV, df=GDF	9	0.760	0.357	0.313	0.579	0.218	0.237
	Fixed	2	0.776	0.568	0.595	1.000	1.000	1.000
	Fixed	3	0.832	0.620	0.594	0.777	0.570	0.597
Mclust	BIC	2	0.776	0.568	0.595	1.000	1.000	1.000
	Fixed	3	0.957	0.904	0.879	0.779	0.572	0.599
Sclust	GCV, df=GDF	9	0.768	0.388	0.344	0.593	0.243	0.264
	Fixed	2	0.776	0.568	0.595	1.000	1.000	1.000
	Fixed	3	0.837	0.630	0.603	0.777	0.569	0.596
HTclust	GCV, df=GDF	4	0.773	0.552	0.578	0.970	0.939	0.94
	Fixed	2	0.776	0.568	0.595	1.000	1.000	1.000
	Fixed	3	0.772	0.558	0.562	0.996	0.991	0.992
PRclust	GCV, df=GDF	3	0.777	0.564	0.589	0.982	0.965	0.968
	Fixed	2	0.776	0.568	0.595	1.000	1.000	1.000
	Fixed	3	0.775	0.562	0.599	0.987	0.974	0.977

Table 5: Results for Fisher’s iris data with 2 or 3 clusters.

3.3 Iris Data

We applied the methods to the popular Fisher’s iris data. There are 4 measurements on the flower, sepal length, sepal width, petal length and petal width, for each observation. There are 50 observations for each of the three iris subtypes. One subtype is well separated from the other two, but the latter two overlap with each other. For this data set, it is debatable whether there are 2 or 3 clusters; for this reason, for any clustering results, we calculated the agreement indices based on the 3 clusters (each corresponding to each iris subtype), and that based on only 2 clusters by combining the latter two overlapping subtypes into one cluster. Since two observations share an equal value on each variable, there are at most $\hat{K} = 149$ clusters.

We standardized the data such that for each variable we had a sample mean 0 and SD=1. We applied the methods to the standardized data ($p = 4$). We used $\nu = 0.4$; we tried a few other values of ν and obtained similar results for GDF. For the K-means, we tried the number of clusters $K = 1, 2, \dots, 30$, each with 20 random starts. For HTclust, we searched 1000 candidate d ’s according to the empirical distribution of the pair-wise distances among the observations. For PRclust, we tried $\lambda_2 \in \{0.1, 0.2, \dots, 2\}$ and $\tau_2 \in \{1.0, 1.1, \dots, 2\}$. The results are shown in Table 5.

For the K-means, in agreement with simulations, the Jump selected perhaps a too large $\hat{K} = 27$, while GCV(GDF) selected $\hat{K} = 9$ perhaps due to the non-spherical shapes of the true clusters (Table 5). Both the K-means with CV1 (or CV2) and Mclust selected $\hat{K} = 2$ and yielded the same clustering results. As for the K-means, GCV(GDF) also selected $\hat{K} = 9$ for Sclust. In comparison, PRclust with

GCV(GDF) yielded $\hat{K} = 3$ clusters with higher agreement indices than those of the K-means, Mclust and Sclust. HTclust selected $\hat{K} = 4$ clusters with the agreement indices less than but close to those of PRclust. We also applied the K-means and Sclust with a fixed $\hat{K} = 2$ or 3, and took the subset of the tuning parameter values yielding 2 or 3 clusters for HTclust and PRclust. It is interesting to note that, with $\hat{K} = 2$, all the methods gave the same results that recovered the two true clusters; however, with $\hat{K} = 3$, the results from PRclust and HTclust were similar, but different from the K-means and Sclust: the K-means and Sclust performed better in terms of the agreement with the 3 true clusters, but less well with the 2 true clusters, than PRclust and HTclust, demonstrating different operating characteristics between the K-means/Sclust and the other two methods. When fixed $\hat{K} = 3$, Mclust gave the best results for $K = 3$, suggesting the advantage of Mclust with overlapping and ellipsoidal clusters.

4. Further Modifications and Comparisons

We explore two well-motivated modifications to our new method, which turn out to be less competitive. Then we demonstrate the performance advantages of our new non-convex penalty over several existing convex penalties.

4.1 Modifications

In PRclust, so far we have fixed $\lambda_1 = 1$, which cannot guarantee $\hat{\theta}_{ij} = \hat{\mu}_i - \hat{\mu}_j$, even approximately (Figure 3a). As an alternative, following Framework 17.1 of Nocedal and Wright (2000), we start the algorithm at $\lambda_1 = 1$, at convergence we increase the value of λ_1 , for example, by doubling its current value, and re-run the algorithm with the parameter estimates from the previous iteration as its starting values; this process is repeated until the convergence when the parameter estimates barely change. As before, we can use the new estimates $\hat{\theta}_{ij}$'s to form clusters. We call this modified method PRclust2. As shown in Figure 3b), for a sufficiently large λ_2 , we'd have all $\theta_{ij} = 0$, leading to all $\hat{\mu}_i$'s (almost) equal in PRclust2; in contrast, no matter how large λ_2 was, we had multiple quite distinct $\hat{\mu}_i$'s in PRclust (Figure 3a). We applied PRclust2 to the earlier examples and obtained the following results: when all the clusters were convex, PRclust2 yielded results very similar to those of PRclust; otherwise, their results were different. Table 6 shows some representative results. It is surprising that PRclust performed better than PRclust2 for simulation Case II with two non-convex clusters. A possible explanation lies in their different estimates of θ_{ij} 's, which are used by both PRclust and PRclust2 to perform clustering. PRclust2 yields $\hat{\theta}_{ij} = \hat{\mu}_i - \hat{\mu}_j$ (approximately) while PRclust does not. PRclust2 forms clusters based on the (approximate) equality of $\hat{\mu}_i$'s, while PRclust clusters two observations i and j together if their $\hat{\mu}_i$ and $\hat{\mu}_j$ are close to each other, say, $\|\hat{\mu}_i - \hat{\mu}_j\|_2 < d_{0,ij}$, where the threshold $d_{0,ij}$ is possibly (i, j) -specific. Hence, PRclust2 seems to be more rigid and greedy in forming clusters than PRclust. Alternatively, we can regard PRclust as an early stopped and thus regularized version of PRclust2; it is well known that early stopping is an effective regularization strategy that avoids over-fitting in neural networks and trees (Hastie et al., 2001, p.326).

PRclust forms a cluster based on a connected component of a graph constructed with $\hat{\theta}_{ij}$'s. More generally, one can apply the spectral clustering of Ng et al. (2002) to either $\hat{\mu}_i$'s or $\hat{\theta}_{ij}$'s obtained in PRclust; we call the resulting method PRclust3 and PRclust4 respectively. We propose using the GCV(GDF) to select both the scale parameter in Sclust and the number of clusters. To reduce computational demand, we manually chose a suitable γ for Sclust. We applied the methods to the

Data	Method	Selection of K	\hat{K}	Rand	aRand	Jaccard
Case I	PRclust2	GCV, df=GDF	2.28	0.980	0.959	0.960
	PRclust3	GCV, df=GDF	2.98	0.923	0.845	0.845
	PRclust4	GCV, df=GDF	3.88	0.876	0.751	0.752
Case II	PRclust2	GCV, df=GDF	2.00	0.498	-0.005	0.329
	PRclust3	GCV, df=GDF	2.00	0.498	-0.005	0.329
	PRclust4	GCV, df=GDF	2.00	0.498	-0.005	0.329
Iris	PRclust2	GCV, df=GDF	3	0.777	0.564	0.589
	PRclust3	GCV, df=GDF	9	0.766	0.392	0.356
	PRclust4	GCV, df=GDF	4	0.777	0.519	0.528

Table 6: Results for modified PRclust for 100 simulated data sets (2 clusters) or the iris data (3 clusters).

data examples; as shown in Table 6, the two methods did not improve over the original PRclust. As a reviewer suggested, alternatively, we may also apply PRclust, not the K-means, to the eigen-vectors in a modified Sclust; however, it will be challenging to develop computationally more efficient methods to simultaneously choose multiple tuning parameters, that is, (γ, k) in Sclust and (λ_2, τ) in PRclust.

4.2 Comparison with Some Convex Fusion Penalties

In contrast to our non-convex gTLP penalty, several authors have studied the use of the L_q -norm-based convex fusion penalties. Pelckmans et al. (2005) proposed using a fusion penalty based on the L_q -norm with the objective function

$$\frac{1}{2} \sum_{i=1}^n \|x_i - \mu_i\|_2^2 + \lambda \sum_{i < j} \|\mu_i - \mu_j\|_q,$$

and proposed an efficient quadratic convex programming-based computing method for $q = 1$. Lindsten et al. (2011) recognized the importance of using a group penalty with $q > 1$, and applied the Matlab CVX package (Grant and Boyd, 2011) to solve the general convex programming problem for the group Lasso penalty with $q = 2$ (Yuan and Lin, 2006). Hocking et al. (2011) exploited the piecewise linearity of the solution paths for $q = 1$ or $q = \infty$, and proposed an efficient algorithm for each of $q = 1, 2$ and ∞ respectively. We call these methods PRclust- L_q . Note that PRclust- L_1 corresponds to our PRclust-Lasso, for which (and our default PRclust-gTLP) however we have proposed a different computing algorithm, the quadratic penalty method. Importantly, due to the use of the convex penalty, the solution path of PRclust- L_q is quite different from that of PRclust-gTLP. Using the Matlab CVX package, we applied PRclust- L_q with $q \in \{1, 2, \infty\}$ to simulation Case I; the results for the first data set are shown in Figure 4. It is clear that the solution path of PRclust- L_1 (Figure 4a) was essentially the same as that of PRclust-Lasso (Figure 3c) (while different computing algorithms were applied). More importantly, overall the solution paths of all three PRclust- L_q were similar to each other, sharing the common feature that the estimated centroids were more and more biased towards the overall mean as the penalty parameter λ increased. This feature of PRclust- L_q makes it

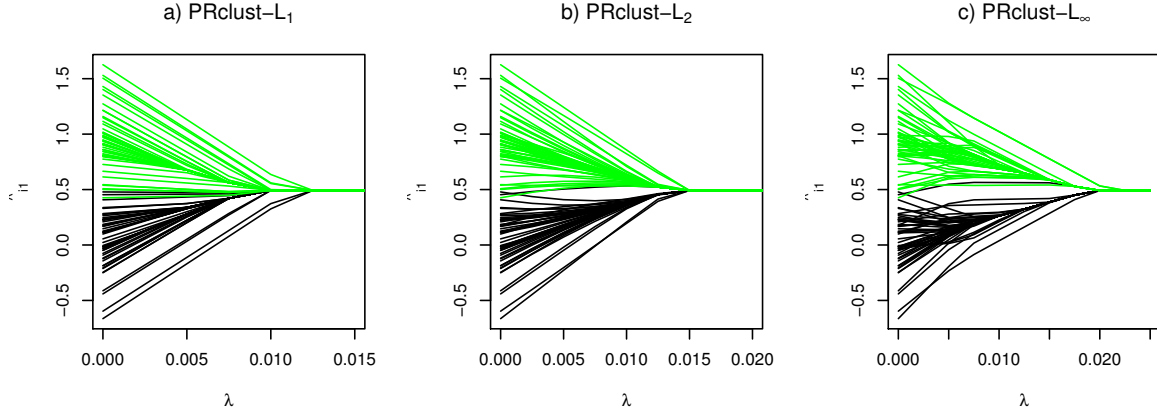


Figure 4: Solution paths of $\hat{\mu}_{i,1}$ for PRclust- L_q with a) $q = 1$, b) $q = 2$ and c) $q = \infty$ for the first simulated data set in Case I.

difficult to correctly select the number of clusters. In fact, both Pelckmans et al. (2005) and Hocking et al. (2011) treated PRclust- L_q as a hierarchical clustering tool; none of the authors discussed the choice of the number of clusters. The issue of an L_q -norm penalty in yielding possibly severely biased estimates is well known in penalized regression, which partially motivated the development of non-convex penalties such as TLP (Shen et al., 2012). In the current context, Lindsten et al. (2011) has recognized the issue of the biased centroid estimates in PRclust- L_q and thus proposed a second stage to re-estimate the centroids after a clustering result is obtained. In contrast, with the use of the non-convex gTLP, the above issues are largely avoided as shown in Figure 3ab).

When we applied the GCV(GDF) to select the number of clusters for PRclust- L_q in simulation Case I, as expected, it performed poorly. Hence, for illustration, we considered an ideal (but not practical) alternative. For any $d_0 \geq 0$, similar to hierarchical clustering, we defined an adjacency matrix $A = (a_{ij})$ with $a_{ij} = I(\|\hat{\mu}_i - \hat{\mu}_j\|_2 \leq d_0)$; any two observations x_i and x_j were assigned to the same cluster if $a_{ij} = 1$. Then for any given $\lambda > 0$ and $d_0 \in \{10^{-1}, 10^{-2}, 10^{-3}, 10^{-4}, 0\}$, we calculated the Rand index for the corresponding PRclust- L_q results and the *true* cluster memberships. We show the results of PRclust- L_q with the values of (λ, d_0) achieving the maximum Rand index, giving an upper bound on the performance of PRclust- L_q with any practical criterion to select the number of clusters. As shown in Table 7, a larger value of q seemed to give better ideal performance of PRclust- L_q ; when compared to PRclust-gTLP (Table 1), none of the three PRclust- L_q methods, even in the ideal case of using the true cluster memberships to select the number of clusters, performed better in selecting the correct number of clusters than PRclust-gTLP with the GCV(GDF) criterion.

5. Discussion

The proposed PRclust clustering bears some similarity to the K-means in terms of the objective in minimizing the sum of squared distances between observations and their cluster centroids, however they differ significantly in their specific formulations, algorithms, and importantly, operating characteristics. Consequently, PRclust can perform much better than the K-means in situations un-

Data	Method	\hat{K}	Rand	aRand	Jaccard
Case I	PRclust- L_1	10.42	0.860	0.719	0.723
	PRclust- L_2	5.56	0.949	0.899	0.900
	PRclust- L_∞	3.06	0.976	0.951	0.952

Table 7: Results for PRclust- L_q with \hat{K} selected by maximizing the Rand index for 100 simulated data sets (2 clusters) in Case I.

suitable or difficult to the K-means, such as in the presence of non-convex clusters, as demonstrated in our simulation Case II (Table 2). Similarly, Mclust does not perform well for non-convex clusters (Table 2), but may have advantages with overlapping and ellipsoidal clusters as for simulation Case I (Table 1) and the iris data (Table 5). There is also some similarity between PRclust and HTclust (or single-linkage hierarchical clustering). Although much simpler, HTclust does not have any mechanism for shrinkage estimation, and in general did not perform better than PRclust in our examples. Between PRclust and spectral clustering, it seems that they are complementary to each other, though it remains challenging to develop competitive model selection criteria for spectral clustering. For example, our results demonstrated the effectiveness of the method of Ng et al. (2002) in selecting the scale parameter γ , but the clustering result also critically depended on the specified k , the number of clusters, for which the GCV(GDF) might not perform well. Although Zelnik-Manor and Perona (2004) have proposed a model selection criterion to self-tune the two parameters γ and $k > 1$, it does not work for $k = 1$; if $k = 1$ is included, the criterion will always select $k = 1$. More generally, model selection is related to kernel learning in spectral clustering (Bach and Jordan, 2006). It is currently an open problem whether the strengths of PRclust and spectral clustering can be combined.

PRclust can be extended in several directions. First, rather than the squared error loss, we can use other loss functions. Corresponding to modifying the K-means to the K-medians, K-midranges or K-modes (Steinley, 2006), we can use an L_1 , L_∞ and L_0 loss function, respectively. Computationally, an efficient coordinate-wise algorithm can be implemented for penalized regression with an L_1 loss (Friedman et al., 2007; Wu and Lange, 2008), but it is unclear how to do so for the other two. K-median clustering is closely related to partitioning-around-centroids (PAM) of Kaufman and Rousseeuw (1990), and is more robust to outliers than is the K-means. A modification of PRclust along this direction may retain this advantage. Second, rather than assuming spherically shaped clusters, as implicitly used by the K-means, we can use a general covariance matrix V with a loss function

$$L(x_i - \mu_i) = \frac{1}{2}(x_i - \mu_i)'V^{-1}(x_i - \mu_i),$$

where V is either given or to be estimated. A non-identity V allows a more general model of ellipsoidal clusters. Alternatively, we can also relax the equal cluster volume assumption and use:

$$L(x_i - \mu_i) = \frac{1}{2}(x_i - \mu_i)'(x_i - \mu_i)/\sigma_i^2,$$

where observation-specific variances σ_i^2 's have to be estimated through grouping pursuit, as for observation-specific means/centroids μ_i 's (Xie et al., 2008). More generally, corresponding to the more general Gaussian mixture model-based clustering (Banfield and Raftery, 1993; McLachlan

and Peel, 2002), we might use

$$L(x_i - \mu_i) = \frac{1}{2}(x_i - \mu_i)'V_i^{-1}(x_i - \mu_i),$$

for a general and observation-specific covariance matrix V_i , though it will be challenging to adopt a suitable grouping strategy to estimate V_i 's effectively. Among others, it might provide a computationally more efficient algorithm than the EM algorithm commonly adopted in mixture model-based clustering (Dempster et al., 1977). Equally, we may accordingly modify the RSS term in GCV so that it will not overly favor spherically shaped clusters. Third, in our current implementation, after parameter estimation, we construct an adjacency matrix and search connected components in the corresponding graph to form clusters. This is a special and simple approach to more general graph-based clustering (Xu and Wunsch, 2005); other more sophisticated approaches may be borrowed or adapted. We implemented a specific combination of PRclust and spectral clustering along with GCV(GDF) for model selection: we first applied PRclust, then used its output as the input to spectral clustering, but it did not show improvement over PRclust. Other options exist; for example, as suggested by a reviewer, it might be more fruitful to replace the K-means in spectral clustering with PRclust. These problems need to be further investigated. Fourth, in the quadratic penalty method, rather than fixing $\lambda_1 = 1$ or allowing $\lambda_1 \rightarrow \infty$, we may want to treat λ_1 as a tuning parameter; a challenge is to develop computationally more efficient methods (e.g., than data perturbation-based GCV estimation) to select multiple tuning parameters. Alternatively, as a reviewer suggested, we may also apply the alternating direction method of multipliers (ADMM) (Boyd et al., 2011), which is closely related to, but perhaps more general and simpler than the quadratic penalty method. Finally, we have not applied the proposed method to high-dimensional data, for which variable selection is necessary. In principle, we may add a penalty into our objective function for variable selection (Pan and Shen, 2007), which again requires a fast method to select more tuning parameters and is worth future investigation.

Perhaps the most interesting idea of our proposal is the view of regarding clustering analysis as a penalized regression problem, blurring the typical line drawn to distinguish clustering (or unsupervised learning) with regression and classification (i.e., supervised learning). This not only opens a door to using various regularization techniques recently developed in the context of penalized regression, such as novel non-convex penalties and algorithms, but also facilitates the use of other model selection techniques. In particular, we find that our proposed regression-based GCV with GDF is promising for the K-means and PRclust (but perhaps not for spectral clustering) in selecting the number of clusters, a hard and interesting problem in itself; since this is not the main point of this paper, we wish to report more on this topic elsewhere.

Acknowledgments

The authors thank Junhui Wang for sharing his R code implementing CV1 and CV2 methods, and thank the action editor and reviewers for their helpful and constructive comments. This research was supported by NSF grants DMS-0906616 and DMS-1207771, and NIH grants R01GM081535, R01HL65462 and R01HL105397.

Appendix A.

We prove Theorem 1 in Section 2.2.

By construction of $S^{(m)}(\mu, \theta)$ and the definition of minimization, for each $m \in N$,

$$\begin{aligned} 0 &\leq S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) = S^{(m+1)}(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) \leq S^{(m)}(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) \\ &\leq S^{(m)}(\hat{\mu}^{(m)}, \hat{\theta}^{(m-1)}) \leq S^{(m)}(\hat{\mu}^{(m-1)}, \hat{\theta}^{(m-1)}) = S(\hat{\mu}^{(m-1)}, \hat{\theta}^{(m-1)}), \end{aligned}$$

implying that $S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)})$ decreases in m . Note that $S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) \geq 0$ for all m . Then it converges, and $S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)})$ must decrease strictly in m before meeting the stopping rule to terminate. Moreover, by construction, $S^{(m+1)}(\mu, \theta)$ has only a finite number of distinctly different functions in m . This implies that $S^{(m+1)}(\hat{\mu}^{(m)}, \hat{\theta}^{(m)}) = S(\hat{\mu}^{(m)}, \hat{\theta}^{(m)})$ has a finite number of different minimizers across all $m \in N$, hence termination must occur finitely.

To show that $(\hat{\mu}^{(m^*)}, \hat{\theta}^{(m^*)})$ is a local minimizer of $S(\mu, \theta)$, we check if it satisfies a local optimality of $S(\mu, \theta)$, defined by regular subdifferentials (Rockafellar and Wets, 2003):

$$[1 + \lambda_1(n-1)]\mu_i - x_i - \lambda_1 \sum_{j>i} (\mu_j + \theta_{ij}) - \lambda_1 \sum_{j<i} (\mu_j - \theta_{ji}) = 0, \quad i = 1, \dots, n, \quad (4)$$

$$-\lambda_1(\mu_i - \mu_j - \theta_{ij}) + \lambda_2 b_{ij} \frac{\theta_{ij}}{\|\theta_{ij}\|_2} = 0, \quad i, j = 1, \dots, n \quad (i < j), \quad (5)$$

where b_{ij} is the regular subdifferential of $\min(\|\theta_{ij}\|_2, \tau)$ at $\|\theta_{ij}\|_2$. Note that $(\hat{\mu}^{(m^*)}, \hat{\theta}^{(m^*)}) = (\hat{\mu}^{(m^*-1)}, \hat{\theta}^{(m^*-1)})$ at termination. Then (4) is satisfied with $(\mu, \theta) = (\hat{\mu}^{(m^*-1)}, \hat{\theta}^{(m^*-1)})$. For (5), we discuss three cases. If $\|\hat{\theta}_{ij}^{(m^*-1)}\|_2 > \tau$, then $\hat{\theta}_{ij}^{(m^*-1)} = \hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}$, implying (5) when $\theta_{ij} = \hat{\theta}_{ij}^{(m^*)}$ because $b_{ij} = 0$. If $0 < \|\hat{\theta}_{ij}^{(m^*)}\|_2 < \tau$ and $\|\hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}\|_2 \geq \frac{\lambda_2}{\lambda_1}$, then

$$\hat{\theta}_{ij}^{(m^*)} = (\|\hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}\|_2 - \frac{\lambda_2}{\lambda_1}) \frac{\hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}}{\|\hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}\|_2},$$

hence that $\|\hat{\theta}_{ij}^{(m^*)}\|_2 = \|\hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}\|_2 - \frac{\lambda_2}{\lambda_1}$. Then (4) is met when $\theta_{ij} = \hat{\theta}_{ij}^{(m^*)}$ because $b_{ij} = 1$. If $0 < \|\hat{\theta}_{ij}^{(m^*)}\|_2 < \tau$ and $\|\hat{\mu}_i^{(m^*)} - \hat{\mu}_j^{(m^*)}\|_2 < \frac{\lambda_2}{\lambda_1}$, then $\|\hat{\theta}_{ij}^{(m^*)}\|_2 = 0$, which is contrary to the fact that $0 < \|\hat{\theta}_{ij}^{(m^*)}\|_2 < \tau$. This completes the proof.

Appendix B.

We prove the equivalence between HTclust and the single-linkage hierarchical clustering (SL-Hclust).

Both the HTclust and SL-Hclust form clusters sequentially and in a finite number of steps; we show that, in each step, the SL-Hclust gives the same clusters as those of HTclust if an appropriate threshold is chosen in the latter. Suppose that the distinct values of $d_{ij} = \|x_i - x_j\|_2$ for all $i \neq j$ are ordered from the smallest to the largest as $d_{(1)}, d_{(2)}, \dots, d_{(m)}$. At the beginning, each observation forms its own cluster in both the HTclust and SL-Hclust. In the next step, the SL-Hclust combines observations according to whether their distances satisfy $d_{ij} \leq d_{(1)}$ to form clusters; in HTclust, if we use a threshold $d_0 = d_{(1)} + \varepsilon$ with a tiny $\varepsilon > 0$, then it results in the same clusters as those of the SL-Hclust. If the clustering results of the two methods are the same in step $k-1 > 0$ and if we use $d_0 = d_{(k)} + \varepsilon$ in HTclust, then it leads to the same clusters as the SL-Hclust in step k . By induction, this completes the proof.

References

- L. An and P. Tao. Solving a class of linearly constrained indefinite quadratic problems by D.C. algorithms. *J. Global Optimization*, 11:253-285, 1997.
- F.R. Bach and M.I. Jordan. Learning spectral clustering, with application to speech separation. *Journal of Machine Learning Research*, 7:1963-2001, 2006.
- J.D. Banfield and A.E. Raftery. Model-based Gaussian and non-Gaussian clustering. *Biometrics*, 49:803-821, 1993.
- S. Boyd, N. Parikh, E. Chu, B. Peleato and J. Eckstein. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends in Machine Learning*, 3(1):1-122, 2011.
- A.P. Dempster, N.M. Laird and D.B. Rubin. Maximum likelihood from incomplete data via the EM algorithm (with discussion). *JRSS-B*, 39:1-38, 1977.
- J. Fan and R. Li. Variable selection via nonconcave penalized likelihood and its oracle properties. *JASA*, 96:1348-1360, 2001.
- C. Fraley and A.E. Raftery. MCLUST Version 3 for R: Normal Mixture Modeling and Model-Based Clustering. Technical Report no. 504, Department of Statistics, University of Washington. 2006.
- J. Friedman, T. Hastie, H. Hofling and R. Tibshirani. Pathwise coordinate optimization. *Ann Appl Statistics*, 2:302-332, 2007.
- G.H. Golub, M. Heath and G. Wahba. Generalized cross-validation as a method for choosing a good ridge parameter. *Technometrics*, 21:215-223, 1979.
- M. Grant and S. Boyd. CVX: Matlab software for disciplined convex programming, version 1.21, 2011. <http://cvxr.com/cvx>
- T. Hastie, R. Tibshirani and J. Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer-Verlag, 2001.
- T. Hocking, A. Joulin, F. Bach and J.-P. Vert. Clusterpath: An Algorithm for Clustering using Convex Fusion Penalties. In L. Getoor and T. Scheffer (Eds.), *Proceedings of the 28th International Conference on Machine Learning (ICML'11)*, p.745-752, 2011.
- L. Hubert and P. Arabie. Comparing partitions. *Journal of Classification*, 2:193-218, 1985.
- P. Jaccard. The distribution of flora in the alpine zone. *New Phytologist.*, 11:37-50, 1912.
- L. Kaufman and P.J. Rousseeuw. Finding Groups in Data: An Introduction to Cluster Analysis. Wiley, New York, 1990.
- F. Lindsten, H. Ohlsson and L. Ljung, L. Clustering using sum-of-norms regularization: With application to particle filter output computation. *2011 IEEE Statistical Signal Processing Workshop (SSP)*, pages 201-204. 2011.

- G.J. McLachlan and D. Peel. *Finite Mixture Model*. New York, John Wiley & Sons, Inc. 2002.
- A. Ng, M. Jordan and Y. Weiss. On spectral clustering: analysis and an algorithm. *NIPS*, 2002.
- J. Nocedal and S.J. Wright. *Numerical Optimization*. Springer, 2000.
- W. Pan and X. Shen. Penalized model-based clustering with application to variable selection. *Journal of Machine Learning Research*, 8:1145-1164, 2007.
- K. Pelckmans, J. De Brabanter, J.A.K. Suykens and B. De Moor. Convex Clustering Shrinkage. Workshop on Statistics and Optimization of Clustering Workshop (PASCAL), London, U.K., Jul. 2005. Available at ftp://ftp.esat.kuleuven.ac.be/pub/SISTA/kpelckma/ccs_pelckmans2005.pdf.
- W.M. Rand. Objective criteria for the evaluation of clustering methods. *JASA*, 66:846-850, 1971.
- R.T. Rockafellar and R.J. Wets. *Variational Analysis*. Springer-Verlag, 2003.
- X. Shen and J. Ye. Adaptive model selection. *JASA*, 97:210-221, 2002.
- X. Shen and H.-C. Huang. Grouping pursuit through a regularization solution surface. *JASA*, 105:727-739, 2010.
- X. Shen, W. Pan and Y. Zhu. Likelihood-based selection and sharp parameter estimation. *JASA*, 107:223-232, 2012.
- D. Steinley. K-means clustering: a half-century synthesis. *British Journal of Mathematical and Statistical Psychology*, 59:1-34, 2006.
- C.A. Sugar and G.M. James. Finding the number of clusters in a data set: An information theoretic approach. *Journal of the American Statistical Association*, 98:750-763, 2003.
- A. Thalamuthu, I. Mukhopadhyay, X. Zheng and G.C. Tseng. Evaluation and comparison of gene clustering methods in microarray analysis. *Bioinformatics*, 22:2405-2412, 2006.
- R. Tibshirani. Regression shrinkage and selection via the lasso. *JRSS-B*, 58:267-288, 1996.
- R. Tibshirani, M. Saunders, S. Rosset, J. Zhu and K. Knight. Sparsity and smoothness via the fused lasso. *JRSS-B*, 67:91-108, 2005.
- R. Tibshirani, G. Walther and T. Hastie. Estimating the number of clusters in a data set via the gap statistic. *JRSS-B*, 63:411-423, 2001.
- P. Tseng. Convergence of block coordinate descent method for nondifferentiable maximization. *J. Opt. Theory Appl.*, 109:474-494, 2001.
- J. Wang. Consistent selection of the number of clusters via crossvalidation. *Biometrika*, 97:893-904, 2010.
- T.T. Wu and K. Lange. Coordinate descent algorithms for Lasso penalized regression. *Ann Appl Statistics*, 2:224-244, 2008.

- B. Xie, W. Pan and X. Shen. Penalized model-based clustering with cluster-specific diagonal covariance matrices and grouped variables. *Electronic Journal of Statistics*, 2:168-212, 2008.
- R. Xu and D. Wunsch. Survey of clustering algorithms. *IEEE Trans Neural Networks*, 16:645-678, 2005.
- J. Ye. On measuring and correcting the effects of data mining and model selection. *JASA*, 93:120-131, 1998.
- M. Yuan and Y. Lin. Model selection and estimation in regression with grouped variables. *JRSS-B*, 68:49-67, 2006.
- L. Zelnik-Manor and P. Perona. Self-tuning spectral clustering. *NIPS*, 2004.